

# Introduction to Database Systems

Shan-Hung Wu  
CS, NTHU

Why do you need  
a database system?

To store data,  
why not just use a file system?

# Advantages of a Database System

- It answers *queries* fast
  - Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011
  - Q2: among a set of employers, increase the salary by 20% for those who have worked longer than 4 years
- Queries (from multiple users) can execute *concurrently* without affecting each other
- It *recovers* from crash
  - No corrupt data after restart

# Data Model and Queries (1/3)

Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011

Step1: structure your data by following the *relational data model*


- Identify *records* (e.g., web pages, authors, etc.) with the same *fields* in your data and place them into respective *tables*

## blog\_pages

blog_id	url	created	author_id	
33981	ms.com/...	2012/10/31	729	← record
33982	apache.org/...	2012/11/15	4412	

## users

user_id	name	balance
729	Steven Sinofsky	10,235
730	Picachu	NULL



# Data Model and Queries (2/3)

Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011

```
CREATE TABLE blog_pages (  
    blog_id INT NOT NULL AUTO_INCREMENT,  
    url VARCHAR(60),  
    created DATETIME,  
    author_id INT);
```

```
INSERT INTO blog_pages (url, created, author_id)  
VALUES ('ms.com/...', 2012/09/18, 729);
```

## blog\_pages

blog_id	url	created	author_id
33981	ms.com/...	2012/10/31	729
33982	apache.org/...	2012/11/15	4412

# Data Model and Queries (3/3)

Q1: among a set of blog pages, find those pages written by Steven Sinofsky after 2011

Step2: issue queries

```
SELECT b.blog_id
FROM blog_pages b, users u
WHERE b.author_id=u.user_id
      AND u.name='Steven Sinofsky'
      AND b.created >= 2011/1/1;
```

# Advantages of a Database System

- It answers *queries* fast
  - Q1: among a set of web pages, find those pages written by Steven Sinofsky after 2011
  - Q2: among a set of employers, increase the salary by 20% for those who have worked longer than 4 years
- Queries (from multiple users) can execute ***concurrently*** without affecting each other
- It ***recovers*** from crash
  - No corrupt data after restart



# Transactions (1/3)

- Each query, by default, is placed in a ***transaction*** (***tx*** for short) automatically

```
BEGIN TRANSACTION;  
    SELECT b.blog_id  
        FROM blog_pages b, users u  
        WHERE b.author_id=u.user_id  
              AND u.name='Steven Sinofsky'  
              AND b.created >= 2011/1/1;  
COMMIT TRANSACTION;
```

# Transactions (2/3)

- You can group multiple queries in a transaction optionally
- For example, Steven wants to donate \$100 to Pikachu:

```
BEGIN TRANSACTION;  
    UPDATE users  
        SET balance=balance-100  
        WHERE name='Steven Sinofsky';  
    UPDATE users  
        SET balance=balance+100  
        WHERE name='Pikachu';  
COMMIT TRANSACTION;
```

# Transactions (3/3)

- A database ensures the **ACID** properties of transactions
- **Atomicity**
  - All operations in a transaction either succeed (transaction commits) or fail (transaction rollback) together
- **Consistency**
  - After/before each transaction (which commits or rollback), your data do not violate any rule you have set
  - E.g., `blog_pages.author_id` must be a valid `users.user_id`
- **Isolation**
  - Multiple transactions can run concurrently, but cannot interfere with each other
- **Durability**
  - Once a transaction commits, any change it made lives in DB permanently (unless overridden by other transactions)

So, why do you need  
a cloud database system?

# Definition



- A **cloud database** is a database designed to run in the cloud
  - Manages data of tremendous applications (called **tenants**)
- Is MySQL a cloud database?
  - I can run MySQL in a Amazon EC2 VM instance
- No

# What's the Difference

- Ideally, in addition to all features provided by a traditional database, a cloud database should ensure **SAE**:
  - **high Scalability**
    - High max. throughput (measured by Tx/Query per second/minute)
  - **high Availability**
    - Stay on all the time, despite of server/network failure
  - **Elasticity**
    - Add/shutdown physical/virtual servers dynamically based on the current workload and/or capability

# The evolving database landscape

451 Research

## Relational

### Analytic

Hadoop Piccolo Teradata Aster Netezza ParAccel SAP Sybase IQ  
 Hadapt Infobright EMC Greenplum IBM InfoSphere  
 HPCC RainStor Teradata Calpont Action VectorWise HP Vertica

## Non-relational

### NoSQL

MarkLogic Castle DataStax Enterprise Acunu  
 Citrusleaf Hypertable  
 Versant BerkeleyDB Cassandra HBase  
 Oracle NoSQL  
 RethinkDB  
 HandlerSocket\* App Engine  
 McObject Riak Redis-to-go  
 SimpleDB  
 Progress LevelDB DynamoDB  
 Redis  
 Membrain  
 Voldemort Couchbase  
 Objectivity Couchbase  
 MongoDB CouchDB  
 Lotus Notes  
 InterSystems

### Graph

Neo4J  
 InfiniteGraph  
 OrientDB  
 DEX  
 NuvolaBase

### Big tables

App Engine

### -as-a-Service

### -as-a-Service

FathomDB  
 Amazon RDS Database.com  
 Postgres Plus Cloud ClearDB  
 Rackspace MySQL Cloud  
 Google Cloud SQL SQL Azure  
 Action Ingres  
 EnterpriseDB  
 SAP Sybase ASE


## NewSQL

### -as-a-Service

StormDB  
 Xeround  
 GenieDB  
 ScaleDB  
 MySQL Cluster  
 ScaleBase  
 NuoDB VoltDB  
 MemSQL JustOneDB  
 Drizzle Akiban Translattice  
 SchoonerSQL  
 Clustrix  
 ScaleArc  
 ParElastic  
 Continuent  
 CodeFutures  
 Clustering/sharding

## Operational

# Why So Complicated?





# Full DB functions + SAE = a goal no one can achieve (currently)

- Even with Oracle 11g + SPARC SuperCluster
  - 30,249,688 TPC-C transactions per minute
  - \$30+ million USDyou loss elasticity

# Primer Goal of This Class

- To guide you through the design trade-offs of existing database systems in the cloud
- **NOT** a comprehensive survey
- **NOT** on how to use the cloud databases

# Goals

- To use a DB system (in only 2 weeks)
- To understand how to *write* a DB system
  - Architecture
  - Trade-offs

# Prerequisites

- Data structure
- Good programming skill
  - OOP (in Java)
  - Multi-threaded programming
  - Project management tools like Git

# Syllabus

- [Here](#)
  - Subject to change
- Mon: lecture
- Thu: labs (TA time)
  - Explain your new assignment
  - Review your pass assignment
- Homework every 2 weeks
  - Not only code, but also reports
  - Summarizing *diff from TA's code* and *your observations*

# Grading


- Homework (x5): 50%
- SQL & Java quiz: 15%
  - *On next Thu (2/29)*
- Midterm exam: 15%
- Final project: 20%
- Q/A Bonus: up to 5%

# Resources

- Text Book
  - Lecture notes
  - Reference links
- Course page
  - <http://www.cs.nthu.edu.tw/~shwu>
- TODO
  - Register your team

# About the Quiz...

- Coverage:
  - Java concurrency
  - How to use a DBMS (SQL language)
- On **2/29**
- Assigned readings:
  - [Java Concurrency Tutorial](#)
  - [SQL Tutorial](#)
  - ~~Chaps 2 and 3 on ER & relational models~~





Questions?

# FAQ (1/2)

- Do I need to write programs in this course?
  - A lot!
  - ***We will give extensive coding assignments***
- Do I need to write code with others?
  - Yes, 1~3 students a team

# FAQ (2/2)

- Do we need to come to the class?
  - No, as long as you can pass
- Is this a light-loading class or heavy-loading class?
  - Should be *heavy* to most students
  - *Reserve time, otherwise you will have high chance to fail*